

3. theano multi-class classification

April 4, 2016

0.1 3. Multi-Class Classification in Theano

The aim of this IPython notebook is to show some features of the Python **Theano** library in the field of machine learning. It has been developed by the LISA group at the *University of Montreal* (see: <http://deeplearning.net/software/theano/>). The notebook also relies on other standard Python libraries such as *numpy*, *pandas* and *matplotlib*.

To exemplify the use of **Theano**, this notebook solves the assignments of the *Machine Learning* MOOC provided by **Coursera** (see: <https://www.coursera.org/learn/machine-learning>) and performed in *Stanford University* by **Andrew Ng** (see: <http://www.andrewng.org/>).

The original MOOC assignments should to be programmed with the **Octave** language (see: <https://www.gnu.org/software/octave/>). The idea with this notebook is to provide Python developpers with interesting examples programmed using **Theano**.

This notebook has been developed using the *Anaconda* Python 3.4 distribution provided by **Continuum Analytics** (see: <https://www.continuum.io/>). It requires the **Jupyter Notebook** (see: <http://jupyter.org/>).

About the author: **Francis Wolinski** has an Engineering Degree From *Ecole des Ponts ParisTech* as well as a MSc. in Artificial Intelligence and a PhD. in Computer Science from *Université Pierre et Marie Curie* (UPMC).

0.1.1 3.1 Logistic Regression

3.1.1 Visualizing the data

```
In [1]: import theano
import numpy as np
import theano.tensor as T
import scipy.io as sio
data = sio.loadmat('data/ex3data1.mat')
X, Y = data['X'], data['y'].T[0]
m = X.shape[0]
sel = X[np.random.permutation(m)[0:100],:] # 100 random digits

In [2]: %matplotlib inline
import matplotlib.pyplot as plt

def display_data(X):
    example_width = int(np.sqrt(X.shape[1]))

    # Compute rows, cols
    m, n = X.shape
    example_height = int(n / example_width)

    # Compute number of items to display
    display_rows = int(np.floor(np.sqrt(m)))
```

```

display_cols = int(np.ceil(m / display_rows))

# Between images padding
pad = 1


# Setup blank display
display_array = - np.ones((pad + display_rows * (example_height + pad),\
                           pad + display_cols * (example_width + pad)))

# Copy each example into a patch on the display array
# dataset contains 20 pixel by 20 pixel grayscale images of the digit
curr_ex = 0
for j in range(display_rows):
    for i in range(display_cols):
        if curr_ex >= m:
            break
        # Get the max value of the patch
        max_val = max(abs(X[curr_ex]))
        display_array[(pad + j * (example_height + pad)):(pad + j * (example_height + pad)) +
                      (pad + i * (example_width + pad)):(pad + i * (example_width + pad)) +
                      np.reshape(X[curr_ex], (example_height, example_width)) / max_val
        curr_ex += 1
    if curr_ex >= m:
        break

plt.figure(figsize=(8,6))
plt.axis('off')
plt.imshow(display_array.T, cmap="Greys_r")

```

In [3]: display_data(sel)



3.1.2 Logistic Regression

```
In [4]: X1 = np.c_[np.ones(m), X] # Add intercept term to X
        X1 = X1.T

        # Choose some alpha and lambda values
        alpha = 0.1
        lambda_ = 0.1

        # to compute theta * theta without first parameter which is not regularized
        I = np.eye(X1.shape[0])
        I[0,0] = 0

        num_labels = 10

        all_theta = np.zeros((num_labels, X1.shape[0]))

        for j in range(num_labels):
            print('Feature: %i' % (j+1))
            # Init Theta and Run Gradient Descent
```

```

t = np.zeros(X1.shape[0])
theta = theano.shared(t,name='theta')

x = T.matrix('x')
y = T.vector('y')

h = 1.0 / (1.0 + T.exp(-T.dot(theta,x)))
r = lambda_ * T.dot(T.dot(I,theta),T.dot(I,theta)) / 2 / m
cost = -T.sum(y * T.log(h) + (1.0 - y) * T.log (1.0 - h))/m + r
grad = T.grad(cost,theta)

train = theano.function([x,y],cost,updates = [(theta,theta-alpha*grad)])

num_iters = 500
for i in range(num_iters):
    costM = train(X1,Y==(j+1))

all_theta[j] = theta.get_value()

# number of positive predictions where best prediction == Y-1
accuracy = np.sum(np.argmax(np.dot(all_theta,X1), axis=0) == (Y-1))/len(Y)
print('\nTrain Accuracy: %f' % accuracy)

Feature: 1
Feature: 2
Feature: 3
Feature: 4
Feature: 5
Feature: 6
Feature: 7
Feature: 8
Feature: 9
Feature: 10

Train Accuracy: 0.885000

In [ ]:

```