

## 3. theano multi-class classification

April 16, 2016

### 0.1 3. Multi-Class Classification in Theano

The aim of this IPython notebook is to show some features of the Python **Theano** library in the field of machine learning. It has been developed by the LISA group at the University of Montreal (see: <http://deeplearning.net/software/theano/>). The notebook also relies on other standard Python libraries such as numpy, pandas and matplotlib.

To exemplify the use of **Theano**, this notebook solves the assignments of the Machine Learning MOOC provided by **Coursera** (see: <https://www.coursera.org/learn/machine-learning>) and performed in Stanford University by **Andrew Ng** (see: <http://www.andrewng.org/>).

The original MOOC assignments should to be programmed with the **Octave** language (see: <https://www.gnu.org/software/octave/>). The idea with this notebook is to provide Python developpers with interesting examples programmed using **Theano**.

This notebook has been developed using the Anaconda Python 3.4 distribution provided by **Continuum Analytics** (see: <https://www.continuum.io/>). It requires the **Jupyter Notebook** (see: <http://jupyter.org/>).

About the author: **Francis Wolinski** has an Engineering Degree From Ecole des Ponts ParisTech as well as a MSc. in Artificial Intelligence and a PhD. in Computer Science from Université Pierre et Marie Curie (UPMC).

#### 0.1.1 3.1 Logistic Regression

##### 3.1.1 Visualizing the data

```
In [1]: import theano
import numpy as np
import theano.tensor as T
from theano.tensor.nnet import sigmoid
import scipy.io as sio
data = sio.loadmat('data/ex3data1.mat')
X, Y = data['X'], data['y'].T[0]
m = X.shape[0]

In [2]: %matplotlib inline
import pandas as pnd
import matplotlib.pyplot as plt

def display_data(X, Y):
    example_width = int(np.sqrt(X.shape[1]))

    # Compute rows, cols
    m, n = X.shape
    example_height = int(n / example_width)

    # Compute number of items to display
```

```

display_rows = int(np.floor(np.sqrt(m)))
display_cols = int(np.ceil(m / display_rows))

# Between images padding
pad = 1

# Setup blank display
display_array = - np.ones((pad + display_rows * (example_height + pad),\
                           pad + display_cols * (example_width + pad)))

# Copy each example into a patch on the display array
# dataset contains 20 pixel by 20 pixel grayscale images of the digit
curr_ex = 0
for j in range(display_rows):
    for i in range(display_cols):
        if curr_ex >= m:
            break
        # Get the max value of the patch
        max_val = max(abs(X[curr_ex]))
        display_array[(pad + j * (example_height + pad)):(pad + j * (example_height + pad)) +
                      (pad + i * (example_width + pad)):(pad + i * (example_width + pad)) +
                      np.reshape(X[curr_ex], (example_height, example_width)) / max_val
        curr_ex += 1
    if curr_ex >= m:
        break

plt.figure(figsize=(8,6))
plt.axis('off')
plt.imshow(display_array.T, cmap="Greys_r")

df = pd.DataFrame(Y.reshape((display_rows, display_cols)).T)
df = df.replace(10,0) # 0 is mapped to 10 in Y
print(df)

```

```


In [3]: rnd = np.random.permutation(m)[0:100]
        sel = X[rnd,:] # 100 random digits
        res = Y[rnd]
        display_data(sel, res)

```

```

0 1 2 3 4 5 6 7 8 9
0 2 8 3 3 4 2 5 2 2 3
1 8 2 5 7 5 2 2 6 7 1
2 9 4 9 2 5 3 4 6 0 4
3 5 5 8 8 3 6 3 7 6 1
4 7 6 6 2 0 9 1 4 7 8
5 6 1 7 7 3 9 1 3 7 3
6 2 8 7 3 7 1 8 1 5 6
7 0 2 6 7 0 1 6 7 4 5
8 2 2 8 2 9 8 3 2 7 6
9 4 5 1 4 9 9 4 4 5 9

```



### 3.1.2 Logistic Regression

```
In [4]: X1 = np.c_[np.ones(m), X] # Add intercept term to X
        X1 = X1.T

        # Choose some alpha and lambda values
        alpha = 0.1
        lambda_ = 0.1

        num_labels = 10

        all_theta = np.zeros((num_labels, X1.shape[0]))

        for j in range(num_labels):
            print('Feature: %i' % (j+1))
            # Init Theta and Run Gradient Descent
            t = np.zeros(X1.shape[0])
            theta = theano.shared(t,name='theta')

            x = T.matrix('x')
```

```

y = T.vector('y')

h = sigmoid(T.dot(theta,x)) # 1.0 / (1.0 + T.exp(-T.dot(theta,x)))
reg = lambda_ * T.sum(theta[1:]**2) / 2 / m # regularization on theta
cost = -T.sum(y * T.log(h) + (1.0 - y) * T.log (1.0 - h))/m + reg
grad = T.grad(cost,theta)

train = theano.function([x,y],cost,updates = [(theta,theta-alpha*grad)])

num_iters = 100
for i in range(num_iters):
    costM = train(X1, Y==(j+1))

all_theta[j] = theta.get_value()

# number of positive predictions where best prediction == Y-1
accuracy = np.sum(np.argmax(np.dot(all_theta,X1), axis=0) == (Y-1))/len(Y)
print('\nTrain Accuracy: %f' % accuracy)

Feature: 1
Feature: 2
Feature: 3
Feature: 4
Feature: 5
Feature: 6
Feature: 7
Feature: 8
Feature: 9
Feature: 10

```

Train Accuracy: 0.841600

## 0.1.2 3.2 Neural Networks

### 3.2.1 Model representation

```

In [5]: # load precomputed network parameters
weights = sio.loadmat('data/ex3weights.mat')
theta1, theta2 = weights['Theta1'], weights['Theta2']

```

### 3.2.2 Feedforward Propagation and Prediction

```

In [6]: # layer sizes
input_layer_size = X1.shape[0]
hidden_layer_size = theta2.shape[1]
output_layer_size = theta2.shape[0]

print('input_layer_size: %i' % input_layer_size)
print('hidden_layer_size: %i' % hidden_layer_size)
print('output_layer_size: %i' % output_layer_size)

x = T.matrix('x')

# first layer
y1 = 1 / (1 + np.exp(-np.dot(theta1,X1)))

```

```

    # second layer
    y2 = np.c_[np.ones(m), y1.T] # Add intercept term to y1
    y2 = y2.T
    y3 = 1 / (1 + np.exp(-np.dot(theta2, y2)))

input_layer_size: 401
hidden_layer_size: 26
output_layer_size: 10

In [7]: # number of positive predictions where best prediction == Y-1
        accuracy = np.sum(np.argmax(y3, axis=0) == (Y-1))/len(Y)
        print('Train Accuracy: %f' % accuracy)

Train Accuracy: 0.975200

```